

```
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.layers import Dense, Activation, Dropout, Conv2D, MaxPooling2D, BatchNormalization
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.metrics import binary_crossentropy
from tensorflow.keras import regularizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Model, Sequential

import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from PIL import Image
from IPython.core.display import display, HTML

# stop annoying tensorflow warning messages

import logging
logging.getLogger("tensorflow").setLevel(logging.ERROR)

from pathlib import PosixPath
import random
import matplotlib.pyplot as plt
%matplotlib inline

from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau

root_path = '/content/drive/MyDrive/archive/training_set/training_set'
data_dir = PosixPath(root_path)

image_count = len(list(data_dir.glob("*/")))
print(f"Image count: {image_count}")

Image count: 8007

class_names = os.listdir(root_path)
class_names

['cats', 'dogs']

class_distribution = [len(os.listdir(root_path + '/' + name)) for name in class_names]
class_distribution

[4001, 4006]

seed = random.randint(0,100)
```

```

random.seed(seed)
seed = random.randint(0,1250)
print(f"Current seed : {seed}")

```

Current seed : 549

```

img_path=r'/content/drive/MyDrive/archive/training_set/training_set/cats/cat.1.jpg'
img=plt.imread(img_path)
print (img.shape)
plt.axis('off')
plt.imshow(img)
plt.show()

```

(280, 300, 3)



```

def preprocess (sdir, trsplit, random_seed):
    filepaths=[]
    labels=[]

    classlist=os.listdir(sdir)

    for klass in classlist:

        classpath=os.path.join(sdir,klass)
        flist=os.listdir(classpath)

        for f in flist:

            fpath=os.path.join(classpath,f)
            filepaths.append(fpath)
            labels.append(klass)

    Fseries=pd.Series(filepaths, name='filepaths')
    Lseries=pd.Series(labels, name='labels')
    df=pd.concat([Fseries, Lseries], axis=1)

    # split df into train_df and test_df
    dsplit=(1-trsplit)
    strat=df['labels']
    train_df, valid_df=train_test_split(df, train_size=trsplit, shuffle=True, random_state=

```

```

    print('train_df length: ', len(train_df), ' valid_df length: ', len(valid_df))
    return train_df, valid_df

sdir=r'/content/drive/MyDrive/archive/training_set/training_set'
train_df, valid_df = preprocess(sdir, .8, seed)

    train_df length:  6405   valid_df length:  1602

filepaths=[]
labels=[]

classlist = os.listdir('/content/drive/MyDrive/archive/test_set/test_set')

for klass in classlist:

    classpath=os.path.join(sdir,klass)
    flist=os.listdir(classpath)

    for f in flist:

        fpath=os.path.join(classpath,f)
        filepaths.append(fpath)
        labels.append(klass)

Fseries=pd.Series(filepaths, name='filepaths')
Lseries=pd.Series(labels, name='labels')
test_df=pd.concat([Fseries, Lseries], axis=1)
print('test_df length: ', len(test_df))

    test_df length:  8007

img_size=(100,100)
channels=3
batch_size=32
img_shape=(img_size[0], img_size[1], channels)
length=len(test_df)
test_batch_size=sorted([int(length/n) for n in range(1,length+1) if length % n ==0 and ler
test_steps=int(length/test_batch_size)
print ( 'test batch size: ',test_batch_size, ' test steps: ', test_steps)

    test batch size:  51   test steps:  157

def scalar(img):
    return img # EfficientNet expects pixels in range 0 to 255 so no scaling is required

trgen=ImageDataGenerator(preprocessing_function=scalar, horizontal_flip=True)
tvgen=ImageDataGenerator(preprocessing_function=scalar)

train_gen=trgen.flow_from_dataframe( train_df, x_col='filepaths', y_col='labels', target_s
                                color_mode='rgb', shuffle=True, batch_size=batch_size)

test_gen=tvgen.flow_from_dataframe( test_df, x_col='filepaths', y_col='labels', target_siz
                                color_mode='rgb', shuffle=False, batch_size=test_batch

```

```
valid_gen=tvgen.flow_from_dataframe( valid_df, x_col='filepaths', y_col='labels', target_size=(256, 256),
                                     color_mode='rgb', shuffle=True, batch_size=batch_size)
classes=list(train_gen.class_indices.keys())
```

```
/usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/dataframe_iterator.
    .format(n_invalid, x_col)
Found 6403 validated image filenames belonging to 2 classes.
Found 8005 validated image filenames belonging to 2 classes.
Found 1602 validated image filenames belonging to 2 classes.
```

```
class_count=len(classes)
print(class_count)
```

```
2
```

```
model_name='EfficientNetB3'
base_model=tf.keras.applications.EfficientNetB3(include_top=False, weights="imagenet", input_shape=(256, 256, 3))
x=base_model.output
x=keras.layers.BatchNormalization(axis=-1, momentum=0.99, epsilon=0.001)(x)
x = Dense(256, kernel_regularizer = regularizers.l2(1e-4), activity_regularizer=regularizers.l2(1e-4),
          bias_regularizer=regularizers.l1(0.006), activation='relu')(x)
x=Dropout(rate=.45, seed=123)(x)
x = Flatten()(x)
output=Dense(class_count, activation='softmax')(x)
model=Model(inputs=base_model.input, outputs=output)
model.compile(Adamax(learning_rate=.001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```
Downloading data from https://storage.googleapis.com/keras-applications/efficientnet43941888/43941136 [=====] - 0s 0us/step
43950080/43941136 [=====] - 0s 0us/step
```

```
early_stopping_monitor = EarlyStopping(monitor='val_accuracy', patience=5, restore_best_weights=True)
```

```
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=5,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=0.00001)
```

```
best_model = ModelCheckpoint('/content/drive/MyDrive/bestmodel.hdf5', monitor='accuracy', save_best_only=True)
best_val_acc = ModelCheckpoint('/content/drive/MyDrive/best_val_acc.hdf5', monitor='val_accuracy', save_best_only=True)
```

```
epochs = 3
```

```
history=model.fit(x=train_gen, epochs=epochs, verbose=1, batch_size = 256, validation_data=(val_gen, val_labels),
                  best_model, best_val_acc,
                  shuffle=True, initial_epoch=0)
```

```
Epoch 1/3
201/201 [=====] - 902s 4s/step - loss: 0.5074 - accuracy: 0.0000
Epoch 2/3
201/201 [=====] - 934s 5s/step - loss: 0.3468 - accuracy: 0.0000
```

Epoch 3/3


201/201 [=====] - 898s 4s/step - loss: 0.2863 - accuracy: 0



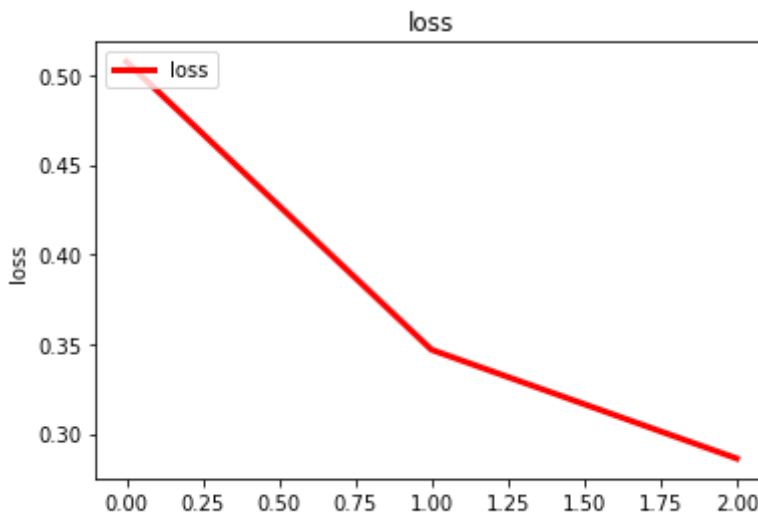
```
def visualization(name,h,color):
    t = h.history[name]
    my_max = max(t)
    my_min = min(t)
    print(f'Name : {name} max : {my_max} min : {my_min}')
    plt.plot(t,color=color,linewidth=3.0)
    plt.title(name)
    plt.ylabel(name)
    plt.xlabel('Epoch')
    plt.legend([name],loc='upper left')
    plt.show()

visualization('accuracy',history,'Blue')
visualization('loss',history,'Red')
visualization('val_accuracy',history,'Green')
visualization('val_loss',history,'Black')
```

Name : accuracy max : 0.98781818151474 min : 0.983601450920105



Name : loss max : 0.5073956847190857 min : 0.28625068068504333



```
import seaborn as sns
from sklearn.metrics import f1_score, confusion_matrix

model.load_weights('/content/drive/MyDrive/bestmodel.hdf5')
accuracy=model.evaluate (test_gen, steps=test_steps)[1]
print(f'accuracy on test set = {accuracy:6.2f}')
```

```
157/157 [=====] - 216s 1s/step - loss: 0.2902 - accuracy: 0
accuracy on test set = 0.99
```



```
from sklearn.metrics import classification_report

predictions = model.predict_generator(test_gen, steps=test_steps)
test_steps_per_epoch = np.math.ceil(test_gen.samples / test_batch_size)
predicted_classes = np.argmax(predictions, axis=1)
true_classes = test_gen.classes
class_labels = list(test_gen.class_indices.keys())
report = classification_report(true_classes, predicted_classes, target_names=class_labels)
print(report)
```


```
cm=confusion_matrix(true_classes,predicted_classes)
print(cm)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: UserWarning: `Model.  
"""Entry point for launching an IPython kernel.
```

	precision	recall	f1-score	support
cats	0.99	0.99	0.99	4000
dogs	0.99	0.99	0.99	4005
accuracy			0.99	8005
macro avg	0.99	0.99	0.99	8005
weighted avg	0.99	0.99	0.99	8005

```
[[3976  24]
 [  44 3961]]
```

```
from sklearn.metrics import ConfusionMatrixDisplay
cm_norm = np.round(cm/np.sum(cm,axis=1).reshape(-1,1),2)
cm_display = ConfusionMatrixDisplay(cm_norm)
cm_display.plot()
plt.show()
```



```
sns.heatmap(cm_norm,cmap='Greens',annot=True,
             cbar_kws={'orientation' : 'vertical','label' : 'Color bar'},
             fmt='.2f'
             )
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
model.load_weights('/content/drive/MyDrive/best_val_acc.hdf5')
accuracy=model.evaluate (test_gen, steps=test_steps)[1]
print(f'accuracy on test set = {accuracy:6.2f}')
```

```
157/157 [=====] - 227s 1s/step - loss: 0.2902 - accuracy: 0.99
accuracy on test set = 0.99
```

```
predictions = model.predict_generator(test_gen, steps=test_steps)
test_steps_per_epoch = np.math.ceil(test_gen.samples / test_batch_size)
predicted_classes = np.argmax(predictions, axis=1)
true_classes = test_gen.classes
class_labels = list(test_gen.class_indices.keys())
report = classification_report(true_classes, predicted_classes, target_names=class_labels)
print(report)
```


```
cm=confusion_matrix(true_classes,predicted_classes)
print(cm)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: UserWarning: `Model.
    """Entry point for launching an IPython kernel.
```


	precision	recall	f1-score	support
cats	0.99	0.99	0.99	4000
dogs	0.99	0.99	0.99	4005
accuracy			0.99	8005
macro avg	0.99	0.99	0.99	8005
weighted avg	0.99	0.99	0.99	8005

```
[[3976  24]
 [  44 3961]]
```

```
cm_norm = np.round(cm/np.sum(cm,axis=1).reshape(-1,1),2)
cm_display = ConfusionMatrixDisplay(cm_norm)
cm_display.plot()
plt.show()
```

```
sns.heatmap(cm_norm,cmap='Greens',annot=True,
            cbar_kws={'orientation' : 'vertical','label' : 'Color bar'},
            fmt='.2f'
            )
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
model.summary()
activation)

block7b_dwconv (DepthwiseConv2D) (None, 4, 4, 2304) 20736 ['block7b_expand
']
block7b_bn (BatchNormalization) (None, 4, 4, 2304) 9216 ['block7b_dwconv
']
block7b_activation (Activation) (None, 4, 4, 2304) 0 ['block7b_bn[0][
']
block7b_se_squeeze (GlobalAveragePooling2D) (None, 2304) 0 ['block7b_activa
']
block7b_se_reshape (Reshape) (None, 1, 1, 2304) 0 ['block7b_se_squ
']
block7b_se_reduce (Conv2D) (None, 1, 1, 96) 221280 ['block7b_se_res
']
block7b_se_expand (Conv2D) (None, 1, 1, 2304) 223488 ['block7b_se_red
']
block7b_se_excite (Multiply) (None, 4, 4, 2304) 0 ['block7b_activa
', 'block7b_se_exp
']
block7b_project_conv (Conv2D) (None, 4, 4, 384) 884736 ['block7b_se_exc
']
block7b_project_bn (BatchNormaliz (None, 4, 4, 384) 1536 ['block7b projec
```

block7b_project_dropout (BatchNormalization)	(None, 4, 4, 384)	1536	['block7b_project_dropout[0][0]']
block7b_drop (Dropout)	(None, 4, 4, 384)	0	['block7b_project_dropout[0][0]']
block7b_add (Add)	(None, 4, 4, 384)	0	['block7b_project_dropout[0][0]']
top_conv (Conv2D)	(None, 4, 4, 1536)	589824	['block7b_add[0][0]']
top_bn (BatchNormalization)	(None, 4, 4, 1536)	6144	['top_conv[0][0]']
top_activation (Activation)	(None, 4, 4, 1536)	0	['top_bn[0][0]']
max_pool (GlobalMaxPooling2D)	(None, 1536)	0	['top_activation[0][0]']
batch_normalization (BatchNormalization)	(None, 1536)	6144	['max_pool[0][0]']
dense (Dense)	(None, 256)	393472	['batch_normalization[0][0]']
dropout (Dropout)	(None, 256)	0	['dense[0][0]']
flatten (Flatten)	(None, 256)	0	['dropout[0][0]']
dense_1 (Dense)	(None, 2)	514	['flatten[0][0]']

```

=====
Total params: 11,183,665
Trainable params: 11,093,290
Non-trainable params: 90,375

```

```
tf.keras.utils.plot_model(model, show_shapes=True)
```

dot: graph is too large for cairo-renderer bitmaps. Scaling by 0.75985 to fit

